

Introduction to Data Structure (Data Management) Lecture 10

Felipe P. Vista IV



Chonbuk National University

- 1 -

Global Frontier College

Reminder

- Everybody, make sure that your name in ZOOM is in the following format:
 - University ID Num Name (no “()”)
 - Ex: 202054321 Juan Dela Cruz
 -
 - Not changing your name to this format
 - you might be marked Absent
 - * Ravshan → absent?



- NoSQL
- JSon and Semi-structured Data

INTRO TO DATA STRUCTURE

NOSQL

(CH 11.1)

Motivation for NoSQL



- Motivated by **Web 2.0** Applications
 - Web 2.0 allow anyone to create and share online information or material
 - Key element is allow people to create, share, collaborate & communicate
 - Hosted services (Google Maps), Web Apps (Google Docs, Flickr), vid sharing sites(YouTube), wikis, blogs, SNS(FB,IG), microblogging(Twitter)

Motivation for NoSQL

- Goal is to scale simple **OLTP**-style applications to millions or even billions of users
- OLTP (OnLine Transaction Processing)
 - capture, store, process data **from transactions in real-time**
 - typical size range from 100MB to 10GB
 - Ex: online banking, purchasing book online, booking ticket, send text message, call center staff view/update customer info

Motivation for NoSQL

- Facebook has 1.79B active users daily (Q2 2020)
 - use often **correlated** in time in each region
 - *correlated* : one thing affects or depends on another
 - **more than 10M** requests/sec if 25% users arrive w/in hour
 - SQL Server would **crash** under this workload
- Users doing both **reads** and **updates**

What is Problem?

- Single server DBMS **too small** for Web data



What is Problem?

- Single server DBMS **too small** for Web data
 - Solution → scale out to multiple servers
 - *scale: resize a device, object or system*
 - *“scale up” or “scale vertically: expanding capability of a machine*
 - *“scale out” or “scale horizontally”: add more machines*



What is Problem?

- Single server DBMS **too small** for Web data
 - Solution → **scale** out to multiple servers
 - *scale: resize a device, object or system*
 - *“scale up” or “scale vertically: expanding capability of a machine*
 - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS




What is Problem?

- Single server DBMS **too small** for Web data
 - Solution → **scale** out to multiple servers
 - *scale: resize a device, object or system*
 - *“scale up” or “scale vertically: expanding capability of a machine*
 - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS
- NoSQL: **reduce** functionality for easier scaling


What is Problem?

- Single server DBMS **too small** for Web data
 - Solution → **scale** out to multiple servers
 - *scale: resize a device, object or system*
 - *“scale up” or “scale vertically: expanding capability of a machine*
 - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS
- NoSQL: **reduce** functionality for easier scaling
 - **simpler** data model
 - **fewer** guarantees

Serverless Architecture



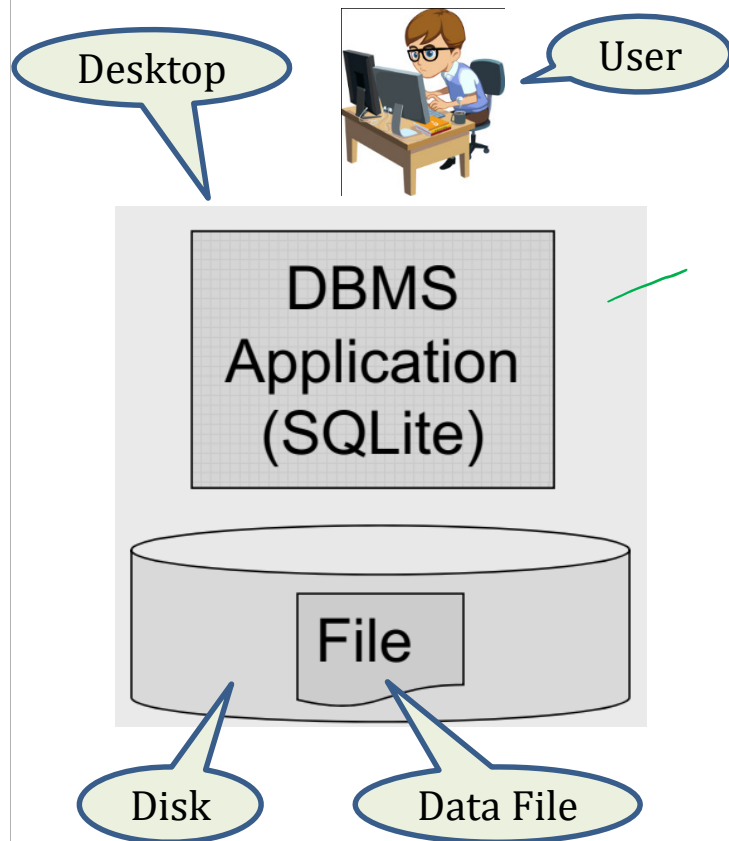
Serverless Architecture



SQLite

- One data file
- One user
- One DBMS application


Serverless Architecture



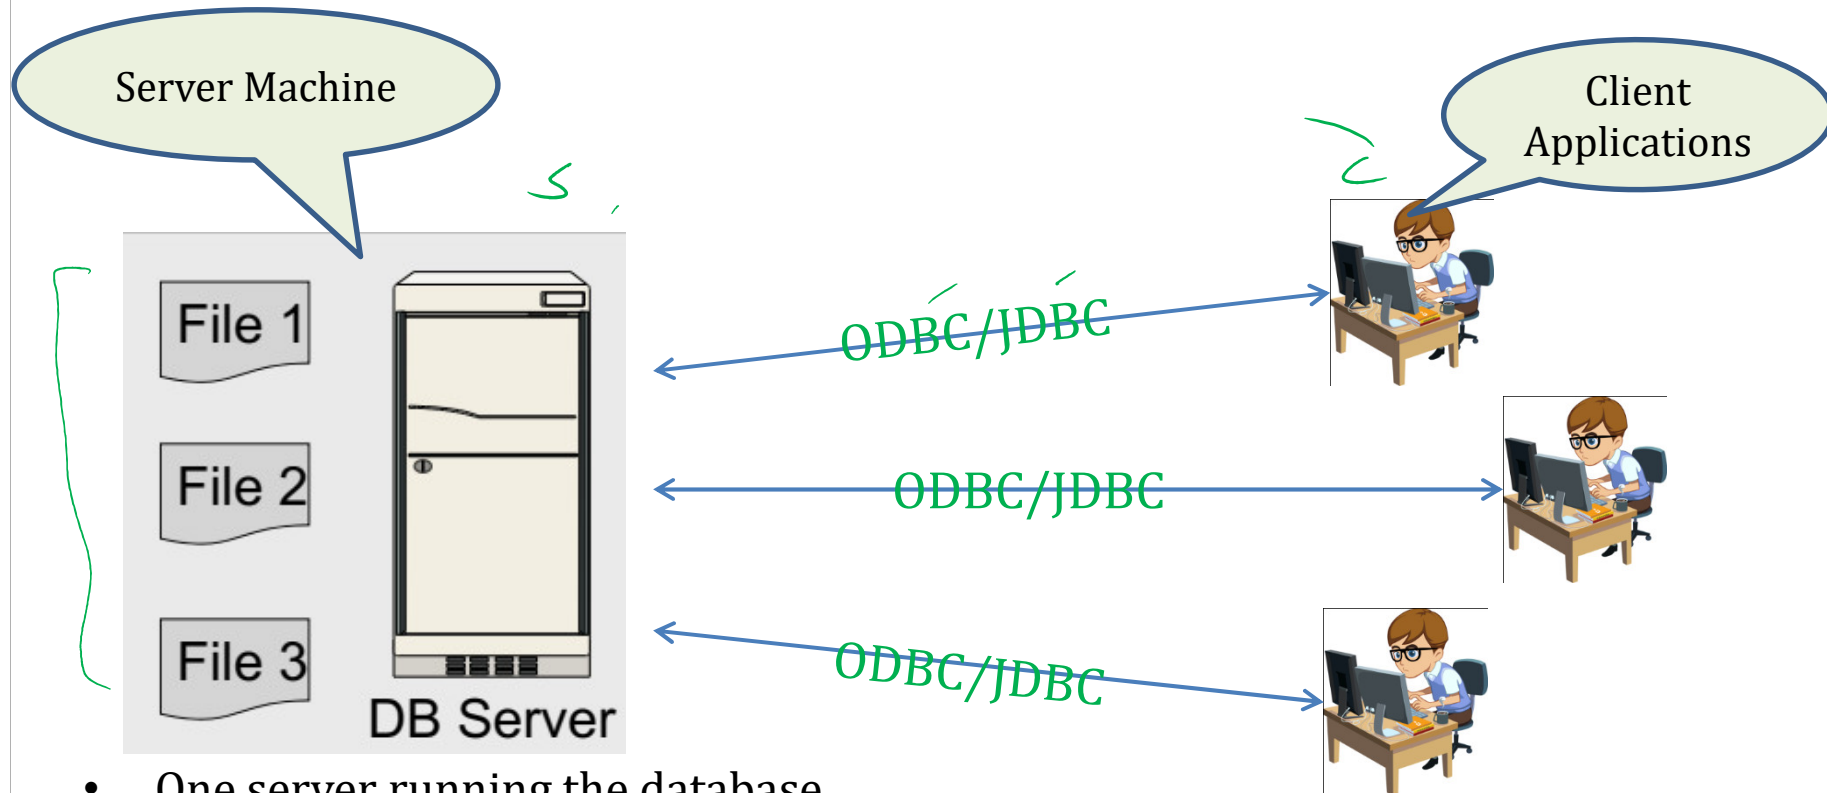
SQLite

- One data file
- One user
- One DBMS application
- Scales well
- But only a **limited** number of scenarios work with such model
- Can be in **browser/ phone**

Client-Server Architecture

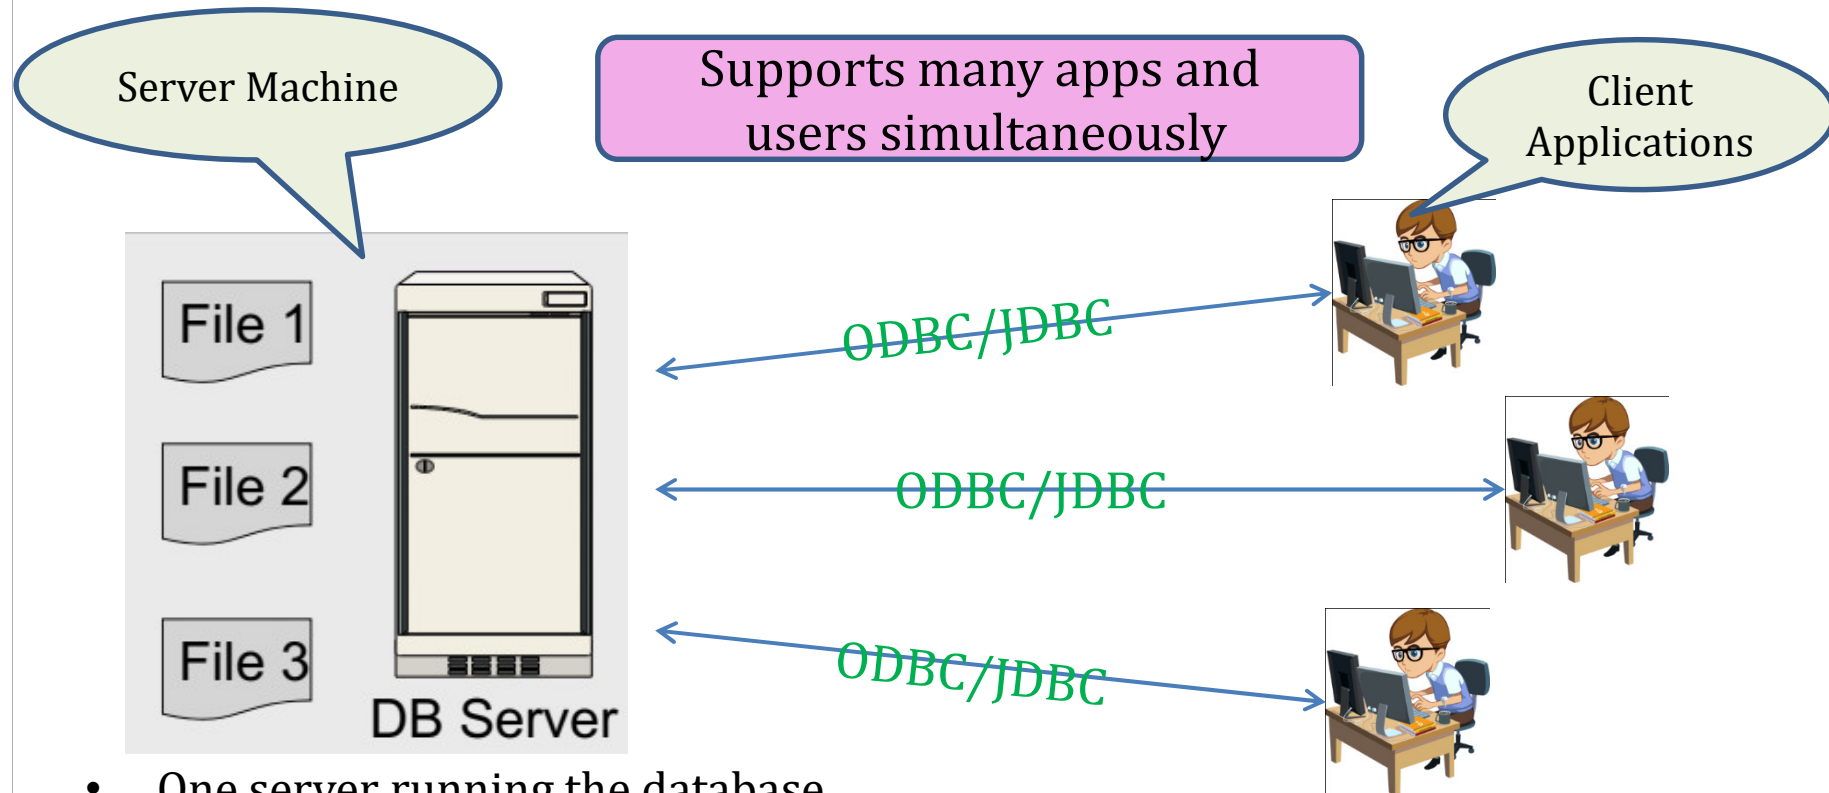


Client-Server Architecture



- One server running the database
- Many clients, connecting via ODBC (Open DB Connectivity / JDBC (Java DB Connectivity

Client-Server Architecture



- One server running the database
- Many clients, connecting via ODBC (Open DB Connectivity / JDBC (Java DB Connectivity

Client-Server

- One *server* runs DBMS (or RDBMS)
 - on own desktop
 - some beefy/powerful system
 - cloud service



Client-Server

- One *server* runs DBMS (or RDBMS)
 - on own desktop
 - some beefy/powerful system
 - cloud service
- Many *clients* run apps and connect to DBMS
 - MS Management Studio (for SQL Server) or
 - pSQL (for postgres)
 - some Java program or C++ program




Client-Server

- One *server* runs DBMS (or RDBMS)
 - on own desktop
 - some beefy/powerful system
 - cloud service
- Many *clients* run apps and connect to DBMS
 - MS Management Studio (for SQL Server) or
 - pSQL (for postgres)
 - some Java program or C++ program
- Clients “*talk*” to server using ODBC/JDBC protocol




3-Tiered Architecture



Web-based Applications

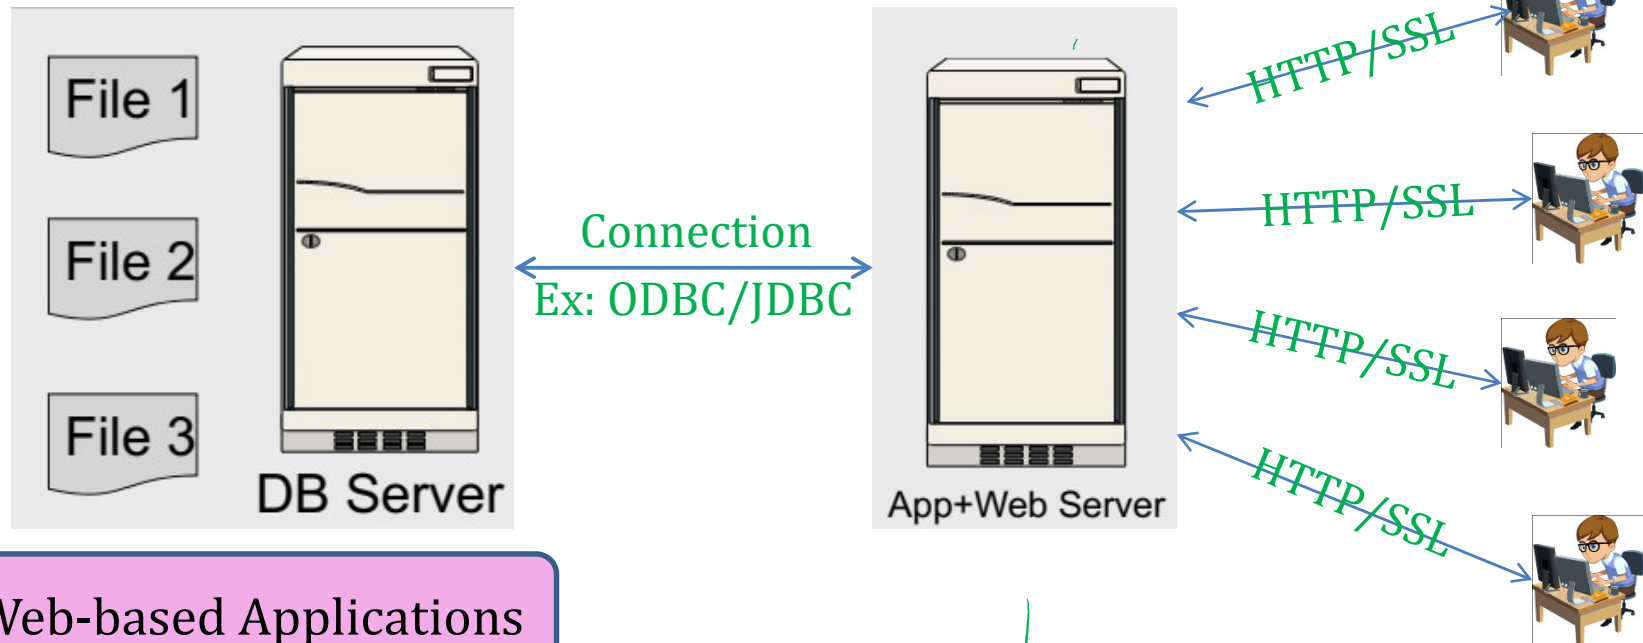


3-Tiered Architecture




Web-based Applications

3-Tiered Architecture




- * HTTP^s = Hyper Text Transfer Protocol
- * SSL = Secure Socket Layer

3-Tiered Architecture



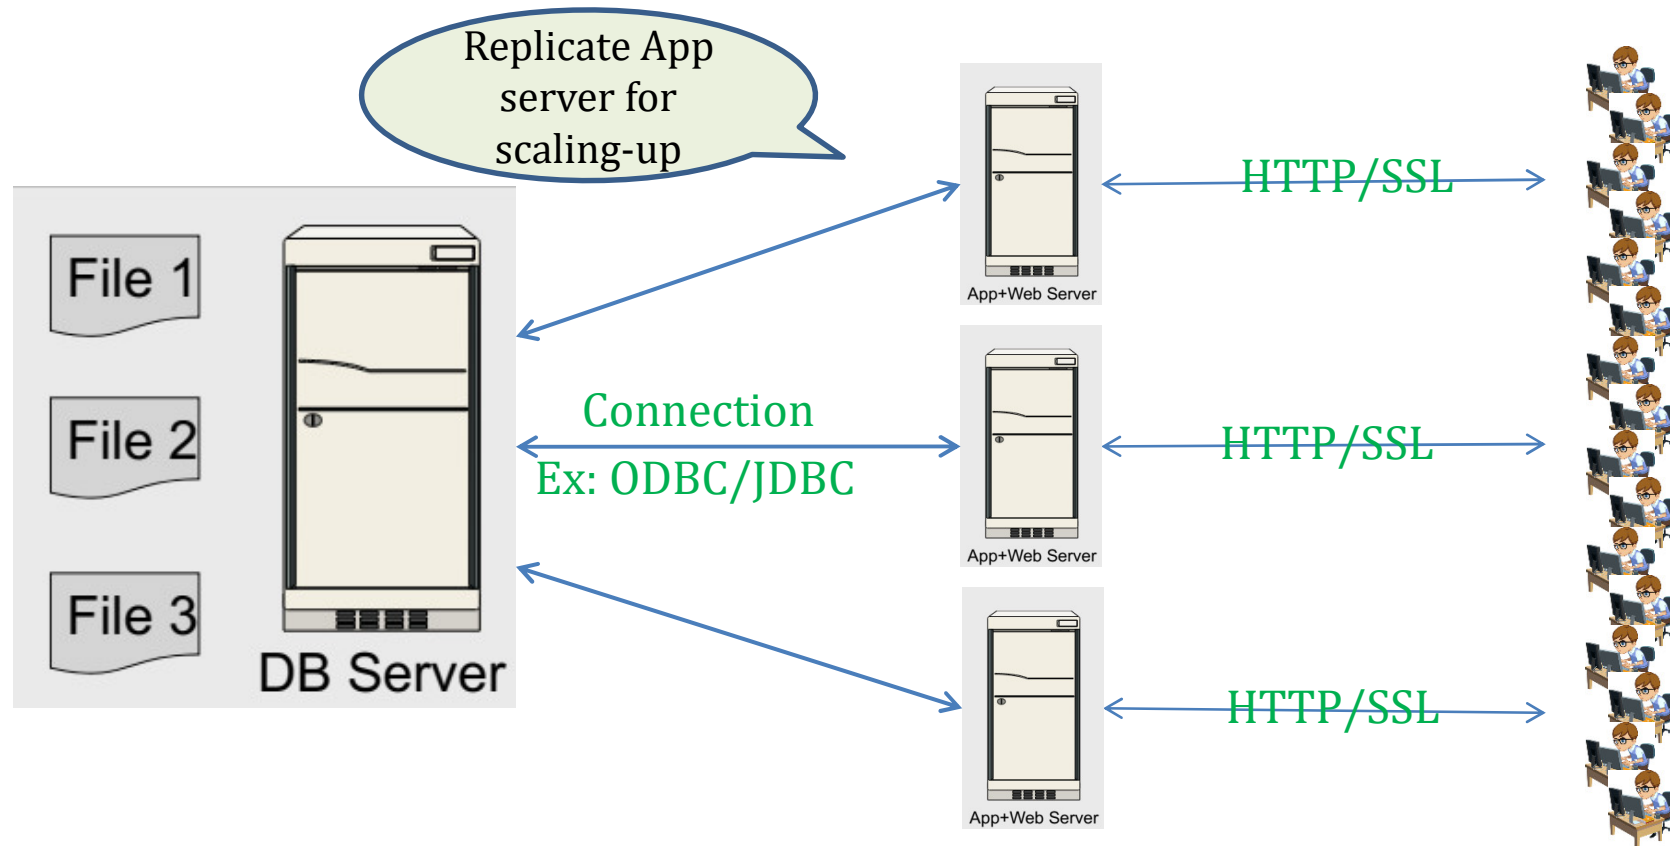
3-Tiered Architecture



* HTTP = Hyper Text Transfer Protocol

* SSL = Secure Socket Layer

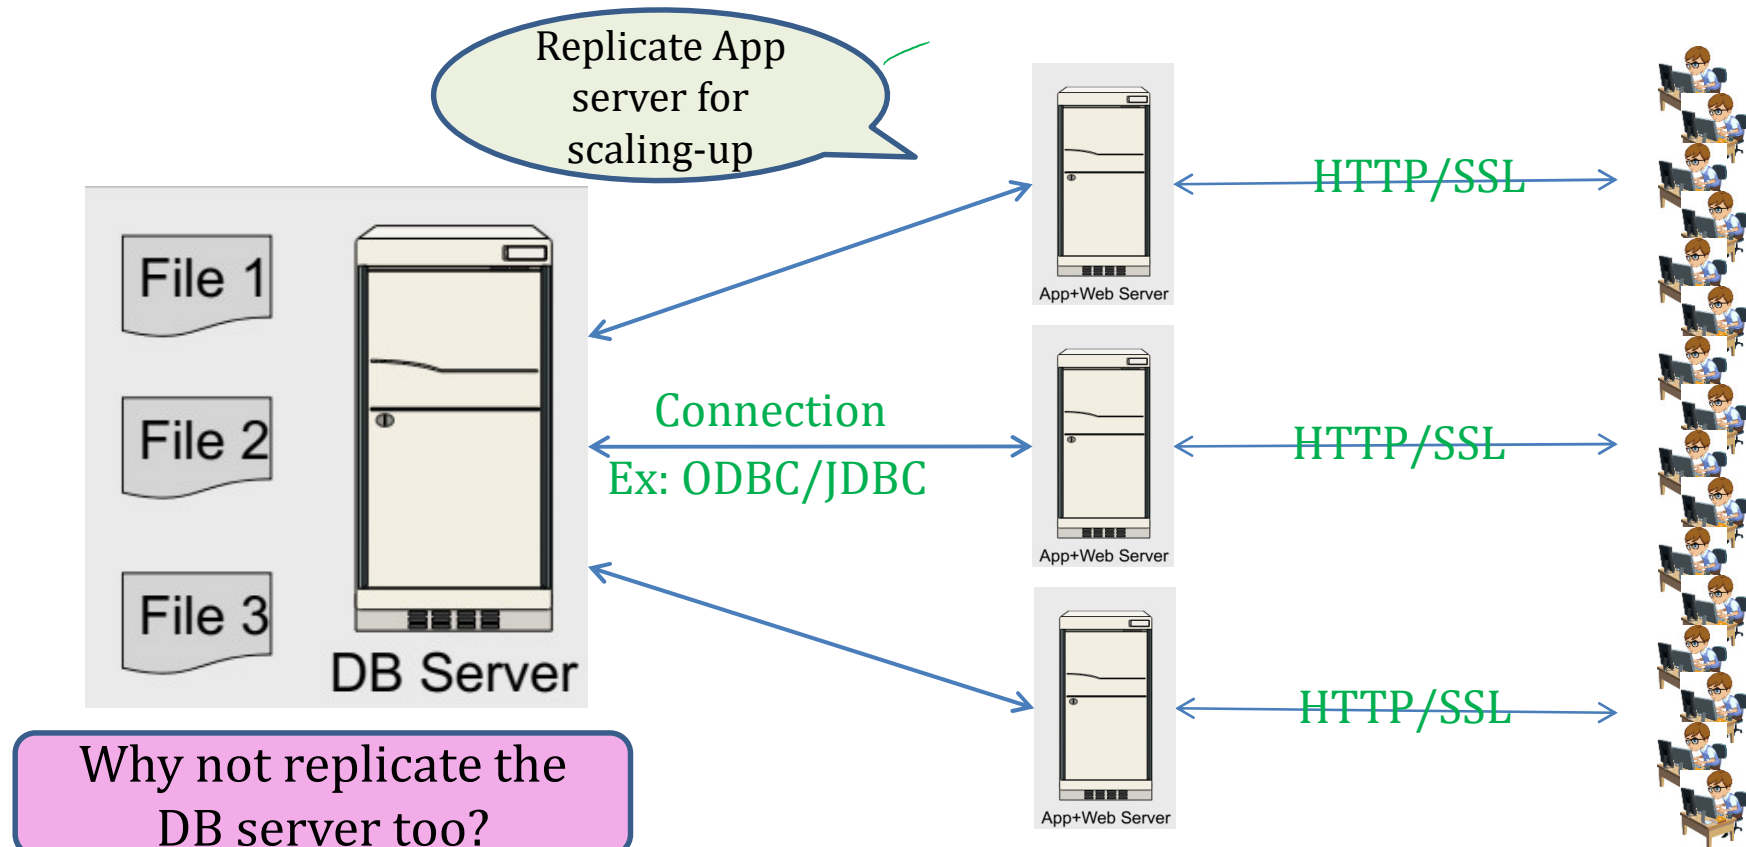
3-Tiered Architecture



* HTTP = Hyper Text Transfer Protocol

* SSL = Secure Socket Layer

3-Tiered Architecture



* HTTP = Hyper Text Transfer Protocol
* SSL = Secure Socket Layer

Replicating the Database


- Much harder because the state must be unique. In other words, database must act as a whole
 - Current DB instance must always be *consistent*
 - Ex: Foreign keys must exist
 - as a result, some *updates* must occur simultaneously

Replicating the Database

- Much harder because the state must be unique. In other words, database must act as a whole
 - Current DB instance must always be *consistent*
 - Ex: Foreign keys must exist
 - as a result, some *updates* must occur simultaneously
- Two basic approach:
 - Scale up by *partitioning*
 - Scale up by *replication*

Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
 - Database could fit in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder



↓
* throughput = amount of material/items
passing through a system or process

Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
 - Database could fit in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder

* throughput = amount of material/items
passing through a system or process



Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
 - Database could fit in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder

