

# Introduction to Data Structure (Data Management) Lecture 10

Felipe P. Vista IV



Chonbuk National University

- 1 -

Global Frontier College

## Reminder

- Everybody, make sure that your name in ZOOM is in the following format:
  - University ID Num Name (no “( )”)
  - Ex: 202054321 Juan Dela Cruz
  - 
  - Not changing your name to this format
    - you might be marked Absent
    - \* → absent?



- NoSQL
- JSon and Semi-structured Data

INTRO TO DATA STRUCTURE

# **NOSQL**

## **(CH 11.1)**

## Motivation for NoSQL



- Motivated by **Web 2.0** Applications
  - Web 2.0 allow anyone to create and share online information or material
  - Key element is allow people to create, share, collaborate & communicate
  - Hosted services (Google Maps), Web Apps (Google Docs, Flickr), vid sharing sites(YouTube), wikis, blogs, SNS(FB,IG), microblogging(Twitter)

## Motivation for NoSQL

- Goal is to scale simple **OLTP**-style applications to millions or even billions of users
- OLTP (OnLine Transaction Processing)
  - capture, store, process data **from transactions in real-time**
  - typical size range from 100MB to 10GB
  - Ex: online banking, purchasing book online, booking ticket, send text message, call center staff view/update customer info



## Motivation for NoSQL

- Facebook has 1.79B active users daily (Q2 2020)
  - use often **correlated** in time in each region
    - *correlated : one thing affects or depends on another*
  - **more than 10M** requests/sec if 25% users arrive w/in hour
  - SQL Server would **crash** under this workload
- Users doing both **reads** and **updates**

## What is Problem?

- Single server DBMS **too small** for Web data





## What is Problem?

- Single server DBMS **too small** for Web data
  - Solution → scale out to multiple servers
    - *scale: resize a device, object or system*
    - *“scale up” or “scale vertically: expanding capability of a machine*
    - *“scale out” or “scale horizontally”: add more machines*



## What is Problem?

- Single server DBMS **too small** for Web data
  - Solution → **scale** out to multiple servers
    - *scale: resize a device, object or system*
    - *“scale up” or “scale vertically: expanding capability of a machine*
    - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS


## What is Problem?

- Single server DBMS **too small** for Web data
  - Solution → **scale** out to multiple servers
    - *scale: resize a device, object or system*
    - *“scale up” or “scale vertically: expanding capability of a machine*
    - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS
- NoSQL: **reduce** functionality for easier scaling


## What is Problem?

- Single server DBMS **too small** for Web data
  - Solution → **scale** out to multiple servers
    - *scale: resize a device, object or system*
    - *“scale up” or “scale vertically: expanding capability of a machine*
    - *“scale out” or “scale horizontally”: add more machines*
- This is hard for **entire** functionality of DBMS
- NoSQL: **reduce** functionality for easier scaling
  - **simpler** data model
  - **fewer** guarantees

# Serverless Architecture



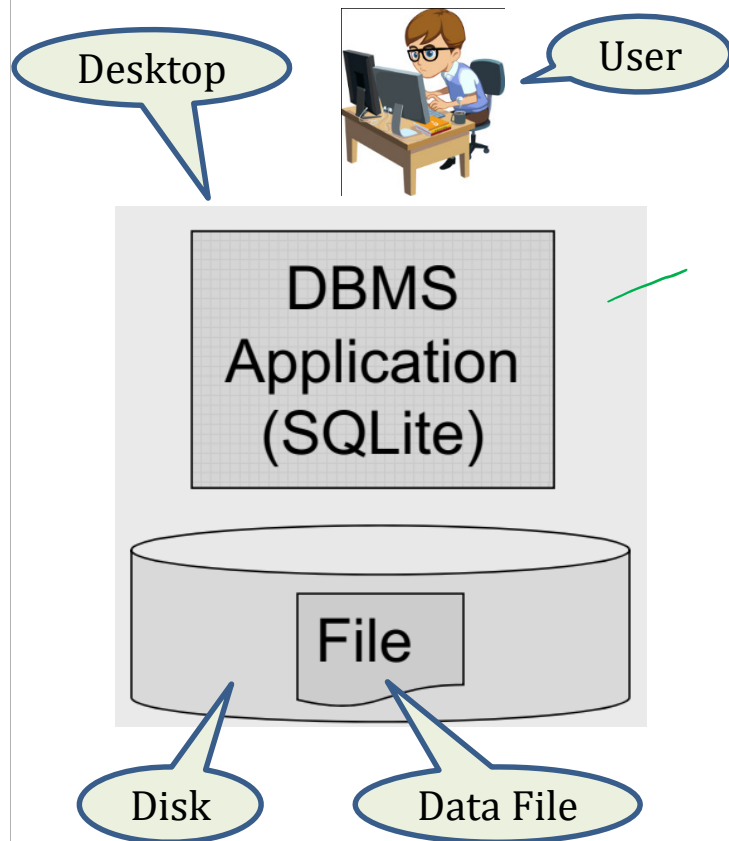
# Serverless Architecture



## SQLite

- One data file
- One user
- One DBMS application


# Serverless Architecture



## SQLite

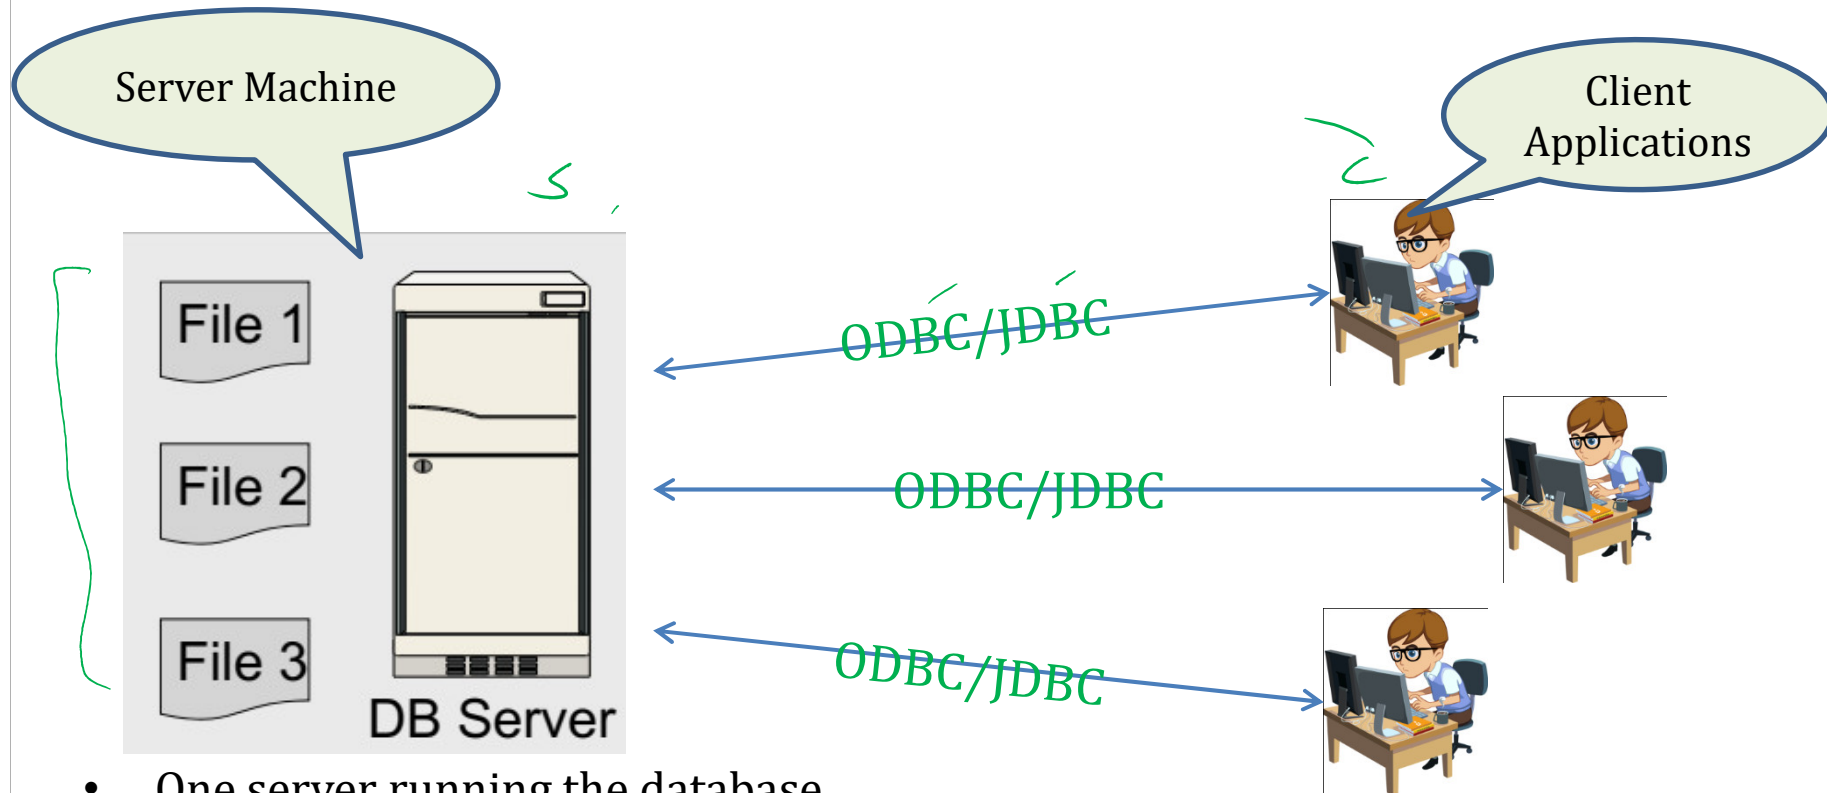
- One data file
- One user
- One DBMS application
- Scales well
- But only a **limited** number of scenarios work with such model
- Can be in **browser/ phone**

# Client-Server Architecture



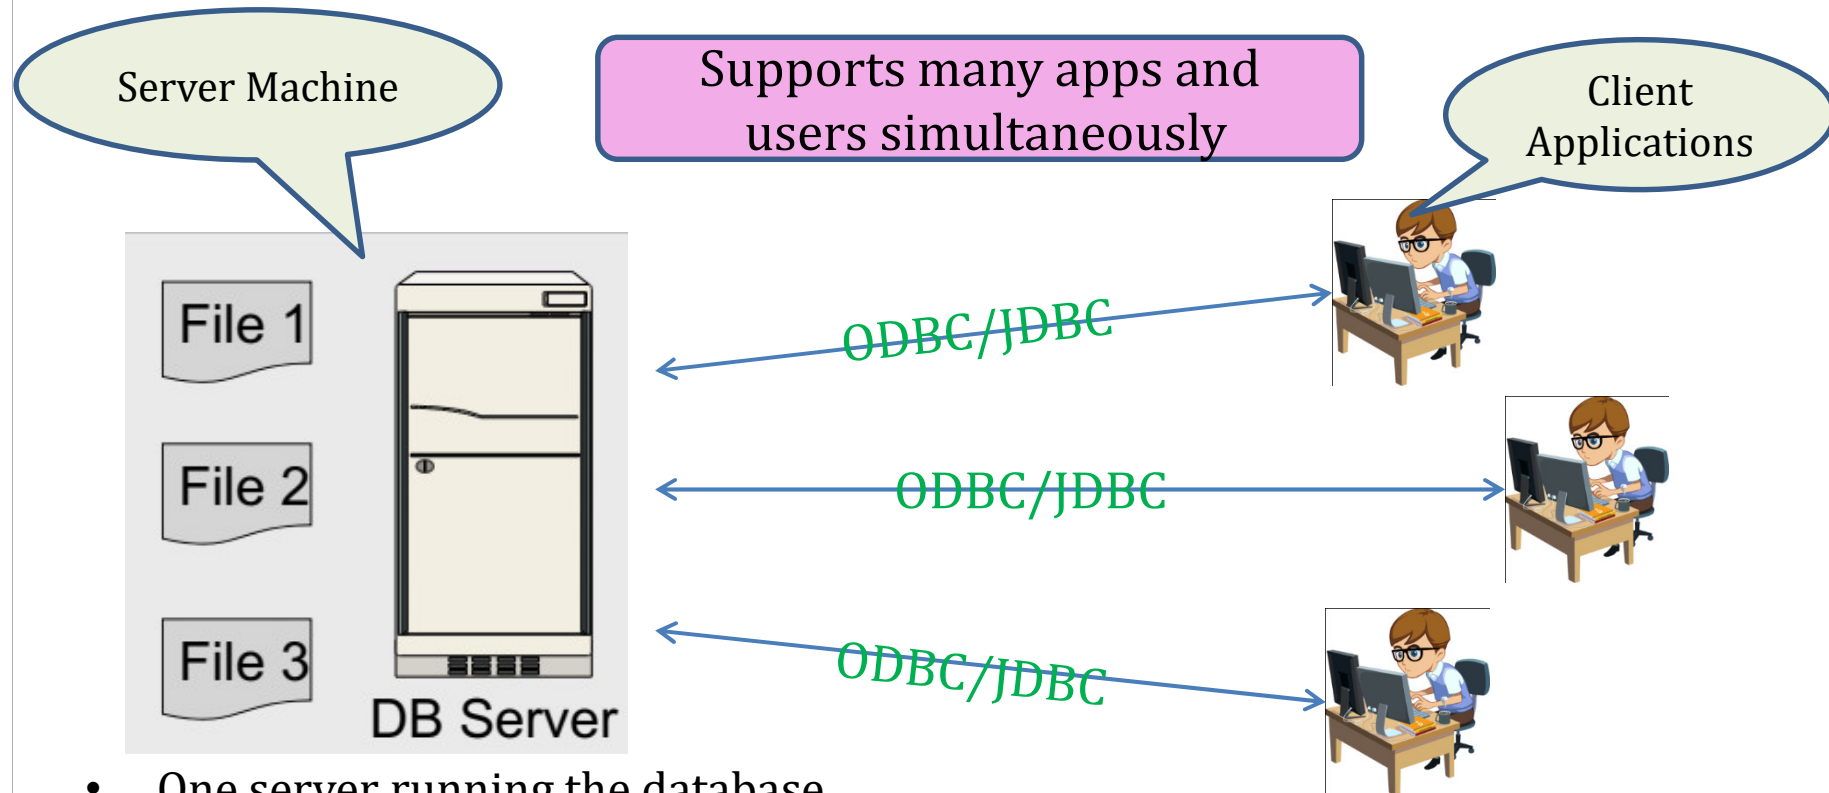


# Client-Server Architecture



- One server running the database
- Many clients, connecting via ODBC (Open DB Connectivity / JDBC (Java DB Connectivity

# Client-Server Architecture



- One server running the database
- Many clients, connecting via ODBC (Open DB Connectivity / JDBC (Java DB Connectivity

## Client-Server

- One *server* runs DBMS (or RDBMS)
  - on own desktop
  - some beefy/powerful system
  - cloud service



## Client-Server

- One *server* runs DBMS (or RDBMS)
  - on own desktop
  - some beefy/powerful system
  - cloud service
- Many *clients* run apps and connect to DBMS
  - MS Management Studio (for SQL Server) or
  - pSQL (for postgres)
  - some Java program or C++ program




## Client-Server

- One *server* runs DBMS (or RDBMS)
  - on own desktop
  - some beefy/powerful system
  - cloud service
- Many *clients* run apps and connect to DBMS
  - MS Management Studio (for SQL Server) or
  - pSQL (for postgres)
  - some Java program or C++ program
- Clients “*talk*” to server using ODBC/JDBC protocol




## 3-Tiered Architecture



Web-based Applications




## 3-Tiered Architecture



Web-based Applications


## 3-Tiered Architecture




- \* HTTP<sup>s</sup> = Hyper Text Transfer Protocol
- \* SSL = Secure Socket Layer



## 3-Tiered Architecture



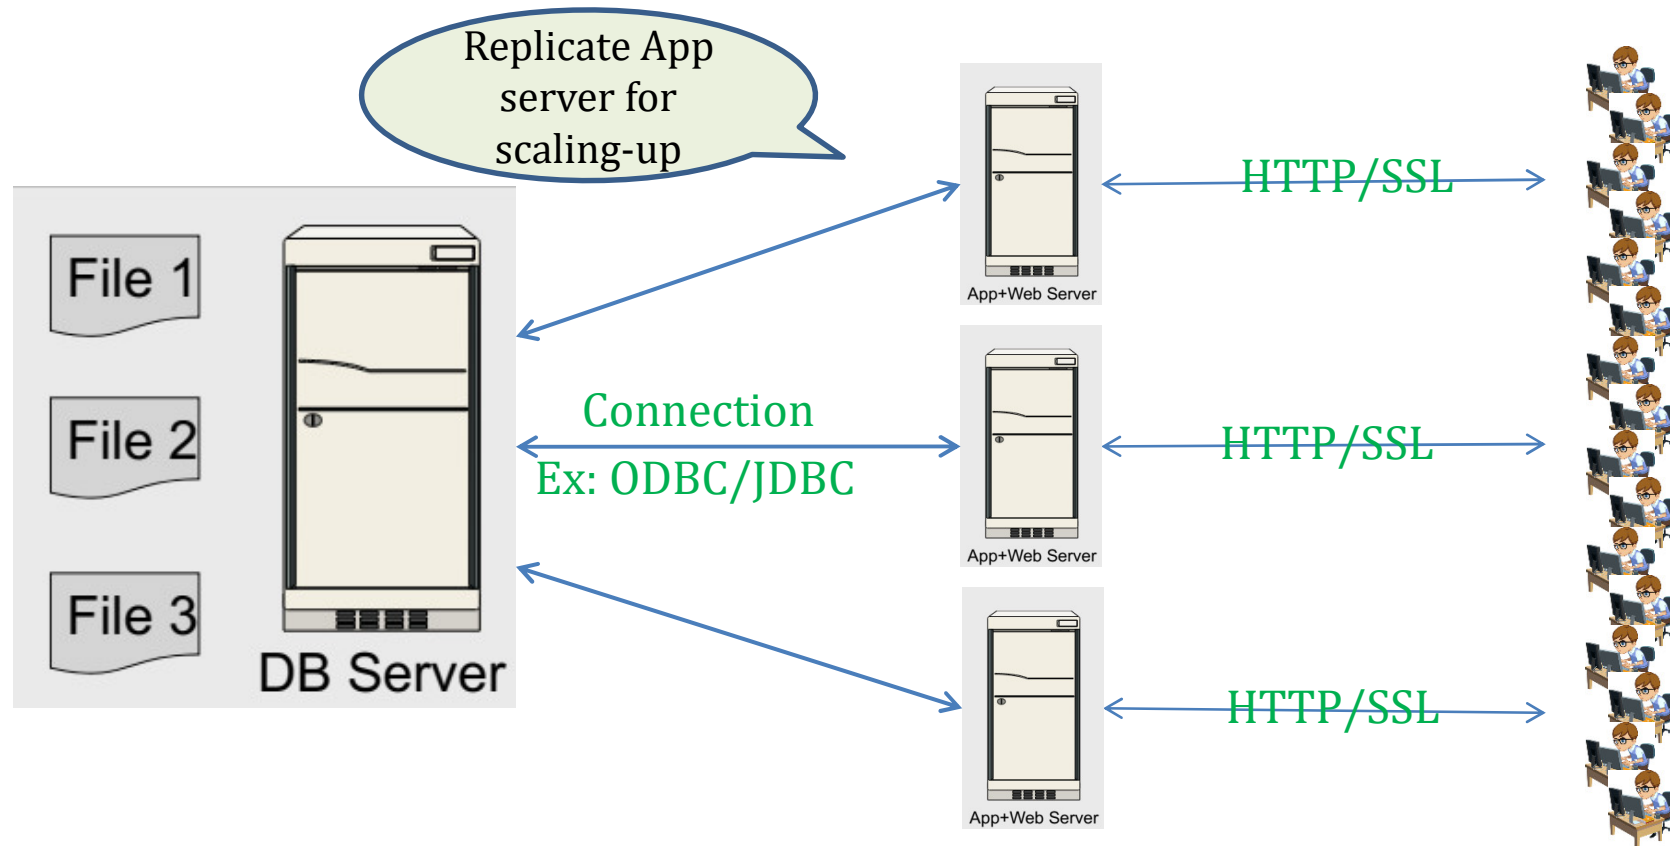
## 3-Tiered Architecture



\* HTTP = Hyper Text Transfer Protocol

\* SSL = Secure Socket Layer

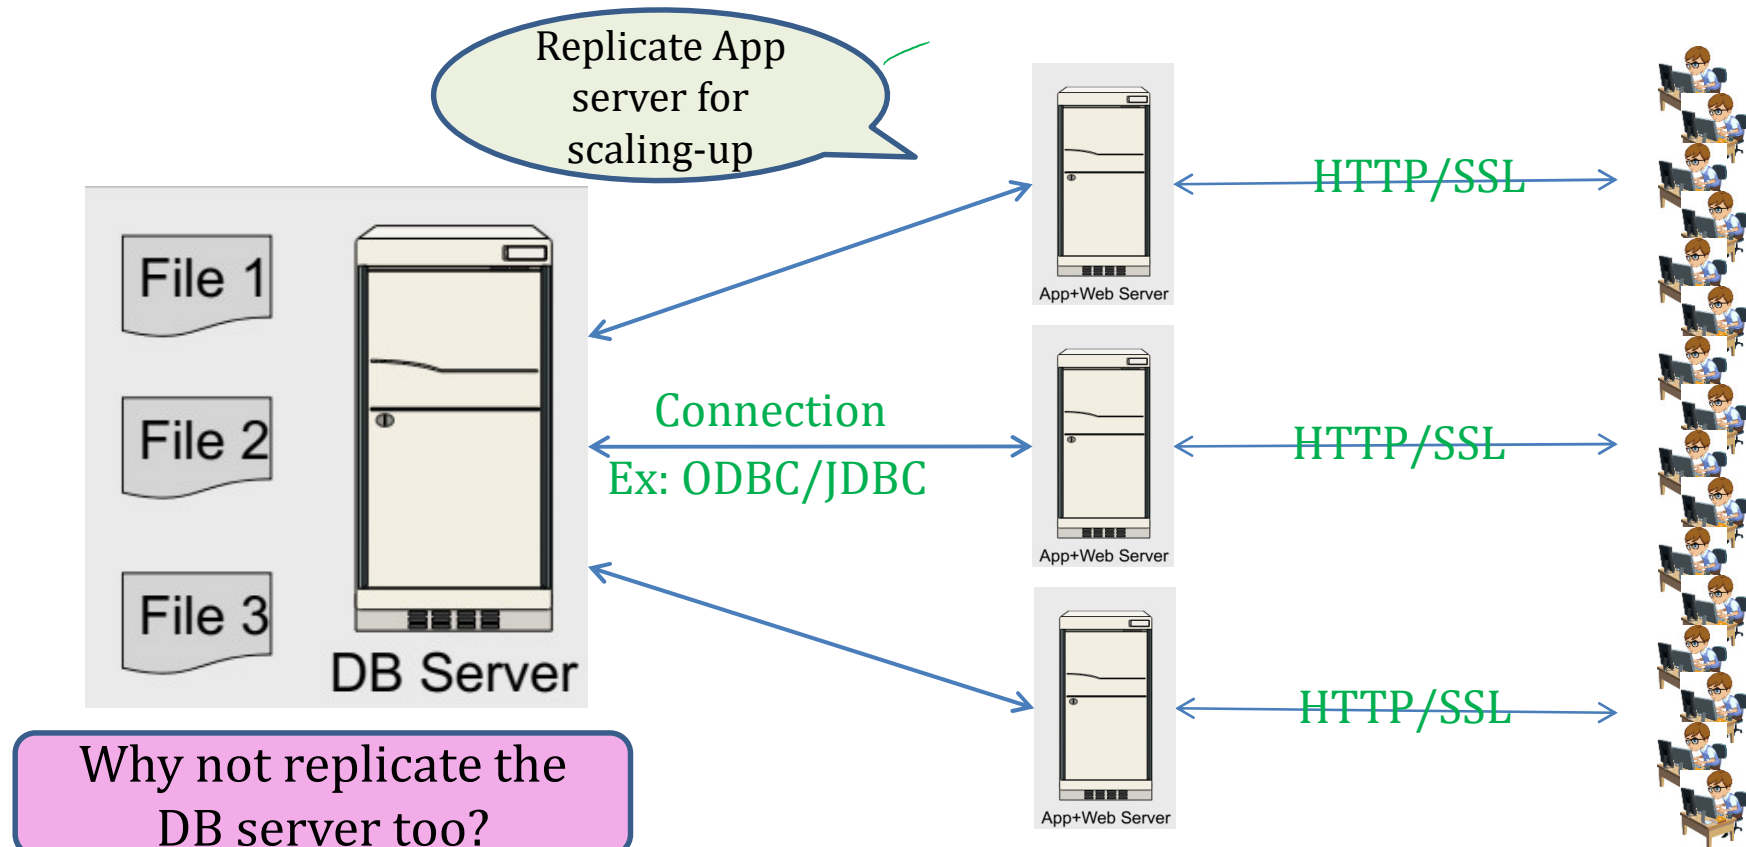
## 3-Tiered Architecture



\* HTTP = Hyper Text Transfer Protocol

\* SSL = Secure Socket Layer

## 3-Tiered Architecture



\* HTTP = Hyper Text Transfer Protocol  
\* SSL = Secure Socket Layer

## Replicating the Database


- Much harder because the state must be unique. In other words, database must act as a whole
  - Current DB instance must always be *consistent*
    - Ex: Foreign keys must exist
    - as a result, some *updates* must occur simultaneously

## Replicating the Database

- Much harder because the state must be unique. In other words, database must act as a whole
  - Current DB instance must always be *consistent*
    - Ex: Foreign keys must exist
    - as a result, some *updates* must occur simultaneously
- Two basic approach:
  - Scale up by *partitioning*
  - Scale up by *replication*

## Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
  - Database could fit in main memory
  - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder



↓  
\* throughput = amount of material/items  
passing through a system or process

## Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
  - Database could fit in main memory
  - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder


\* throughput = amount of material/items  
passing through a system or process





## Scale Through Partitioning

- **Partition** the DB across many machines in a cluster
  - Database could fit in main memory
  - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder



## Scale Through Replication

- Create **multiple copies** of each database partition
- Spread queries across these replicas
- Can increase throughput and lower latency
- Can also improve fault-tolerance
- Easy for reads but writes become harder

\* latency = delay before transfer of data starts after instruction for its transfer


\* fault-tolerance = ability of system to continue operating w/o interruption when one or more components fail

## Scale Through Replication

- Create **multiple copies** of each database partition
- Spread queries across these replicas
- Can increase throughput and lower latency
- Can also improve fault-tolerance
- Easy for reads but writes become harder

\* latency = delay before transfer of data starts after instruction for its transfer

\* fault-tolerance = ability of system to continue operating w/o interruption when one or more components fail



## NoSQL Data Models

Taxonomy based on data models

- **Key-value stores**
  - Ex.: Project Voldemort(LinkedIn), Memcached

\* taxonomy = science of categorization or classification of things  
based on a predetermined system



## NoSQL Data Models

Taxonomy based on data models

- **Key-value stores**
  - Ex.: Project Voldemort(LinkedIn), Memcached
- **Document stores**
  - Ex.: SimpleDB, CouchDB, MongoDB

\* taxonomy = science of categorization or classification of things  
based on a predetermined system



## NoSQL Data Models

Taxonomy based on data models

- **Key-value stores**
  - Ex.: Project Voldemort(LinkedIn), Memcached
- **Document stores**
  - Ex.: SimpleDB, CouchDB, MongoDB
- **Extensible Record stores**
  - Ex.: Hbase, Cassandra, PNUTS

\* taxonomy = science of categorization or classification of things  
based on a predetermined system



# NoSQL Data Models

Taxonomy based on data models

## Key-value stores

- Ex.: Project Voldemort(LinkedIn), Memcached

- Document stores

- Ex.: SimpleDB, CouchDB, MongoDB

- Extensible Record stores

- Ex.: Hbase, Cassandra, PNUTS

\* taxonomy = science of categorization or classification of things  
based on a predetermined system



## Key-Value Stores Features

- **Data Model:** (key, value) pairs
  - Key = string/integer, unique for entire data
  - Value = can be anything (very complex object)



## Key-Value Stores Features

- **Data Model:** (key, value) pairs
  - Key = string/integer, unique for entire data
  - Value = can be anything (very complex object)
- **Operations**
  - Get(key), Put(key, value)
  - Operations on value not supported —

## Key-Value Stores Features

- **Data Model:** (key, value) pairs
  - Key = string/integer, unique for entire data
  - Value = can be anything (very complex object)
- **Operations**
  - Get(key), Put(key, value)
  - Operations on value not supported
- **Distribution/ Partitioning**
  - No replication: key  $k$  is stored at server  $h(k)$
  - 3-way replication: key is stored at  $h1(k)$ ,  $h2(k)$ ,  $h3(k)$

How does  $get(k)$  work? How does  $put(k,v)$  work?



```
- Flights(fid, date, carrier,  
          flight_num, origin, dest, ...)  
- Carriers(cid, name)
```

## Example

How would you represent the **Flights** data as (key, value) pairs

```
Flights(fid, date, carrier,  
        flight_num, origin, dest, ...)  
Carriers(cid, name)
```



## Example

How would you represent the **Flights** data as (key, value) pairs

*date # carrier # flight\_num # ...*

- Option 1: key=fid, value=entire flight record

```
Flights(fid, date, carrier,  
        flight_num, origin, dest, ...)  
Carriers(cid, name)
```

## Example

How would you represent the **Flights** data as (key, value) pairs

- Option 1: key=fid, value=entire flight record

*datetime {oldmmyyhhmmss}*

- Option 2: key=date, value=all flights that day

```
Flights(fid, date, carrier,  
        flight_num, origin, dest, ...)  
Carriers(cid, name)
```

## Example

How would you represent the **Flights** data as (key, value) pairs

- Option 1: key=fid, value=entire flight record —
- Option 2: key=date, value=all flights that day ~
- Option 3: key=(<sup>inc hwi</sup>origin, dest), value=all flights between —

How does query processing work?



# Key-Value Stores Internals

- Data remains in main memory
  - One implementation: [distributed hash table](#)



## Key-Value Stores Internals

- Data remains in main memory
  - One implementation: [distributed hash table](#)
- Most systems also offer a [persistence](#) option


\* persistence = data survives after process it was created has ended





## Key-Value Stores Internals

- Data remains in main memory
  - One implementation: **distributed hash table**
- Most systems also offer a **persistence** option
- Others use **replication** to provide fault-tolerance
  - **Asynchronous** replication: copy data to the replica after the data is already written to the primary storage
  - **Synchronous** replication: write data to primary storage and the replica simultaneously
  - **Tunable** consistency: read/write one replica or majority





\* persistence = data survives after process it was created has ended

\* replica = exact copy of a database or other data store

## Key-Value Stores Internals

- Some offer **transactions**, others do not
  - Multi-version concurrency control or locking
- No** secondary indices



## Data Models

### Taxonomy based on data models

- Key-value stores

- Ex.: Project Voldemort(LinkedIn), Memcached



### Document stores

- Ex.: SimpleDB, CouchDB, MongoDB

- Extensible Record stores

- Ex.: Hbase, Cassandra, PNUTS

\* taxonomy = science of categorization or classification of things  
based on a predetermined system



## Document Stores Features

- **Data Model:** (key, document) pairs
  - Key = string/integer, unique for entire data
  - Document = JSON or XML

## Document Stores Features

- **Data Model:** (key, document) pairs
  - Key = string/integer, unique for entire data
  - Document = JSON or XML
- **Operations**
  - Get/Put document by key  $\rightarrow$   $Get(k, d)$   $P(k, d)$
  - Limited, non-standard query language on JSON

K V  
 $\hookrightarrow Get(k), P(k, v)$

## Document Stores Features


- **Data Model:** (key, document) pairs
  - Key = string/integer, unique for entire data
  - Document = JSON or XML
- **Operations**
  - Get/Put document by key
  - Limited, non-standard query language on JSON
- **Distribution/ Partitioning**
  - Entire documents, as for key/value pairs

Will discuss JSon next time



## Data Models

Taxonomy based on data models

- Key-value stores
  - Ex.: Project Voldemort, Memcached
- Document stores
  - Ex.: SimpleDB, CouchDB, MongoDB
-  Extensible Record stores
  - Ex.: Hbase, Cassandra, PNUTS

\* taxonomy = science of categorization or classification of things based on a predetermined system



## Extensible Record Stores

- Based on Google's **BigTable** (data storage system)
  - HBase is an open source implementation of BigTable





## Extensible Record Stores

- Based on Google's **BigTable** (data storage system)
  - HBase is an open source implementation of BigTable
- Data model is **rows and columns**
  - Can add both new rows and new columns

Handwritten diagram illustrating a BigTable structure:


	C1	C2	C3
R1	-	-	-
R2	+	-	-

Long Schema

C1: INT, PK  
C2: varchar(2)

## Extensible Record Stores

- Based on Google's **BigTable** (data storage system)
  - HBase is an open source implementation of BigTable
- Data model is **rows and columns**
  - Can add both new rows and new columns
- Scalability by **splitting rows & columns over nodes**
  - Rows partitioned through hashing on primary key
  - Columns of a table are distributed over multiple nodes using “column groups”



## NoSQL Summary

- Simple data model with weaker guarantees
- But they scale as far as needed



## NoSQL Summary

- Simple data model with weaker guarantees
- But they scale as far as needed
- Meanwhile...  
SQL systems continue to **improve**

## Recent SQL Progress

- Modern systems need to store data across the globe
  - Individual data centers go offline
  - Need servers close to users to be efficient



## Recent SQL Progress

- Modern systems need to store data across the globe
  - Individual data centers go offline
  - Need servers close to users to be efficient
- Speed of light is fundamental limit
  - 200+ms latency (across US) visible to users
  - 40ms latency (South Korea 4G networks)



## Recent SQL Progress

- Modern systems need to store data across the globe
  - Individual data centers go offline
  - Need servers close to users to be efficient
- Speed of light is fundamental limit
  - 200+ms latency (across US) visible to users
  - 40ms latency (South Korea 4G networks)
- Systems must weaken guarantees



## Recent SQL Progress

- Modern systems need to store data across the globe
  - Individual data centers go offline
  - Need servers close to users to be efficient
- Speed of light is fundamental limit
  - 200+ms latency (across US) visible to users
  - 40ms latency (South Korea 4G networks)
- Systems must weaken guarantees
- Google Spanner (support SQL)
  - Write data over whole globe (a bit slowly)
  - Reads occur slightly in the past



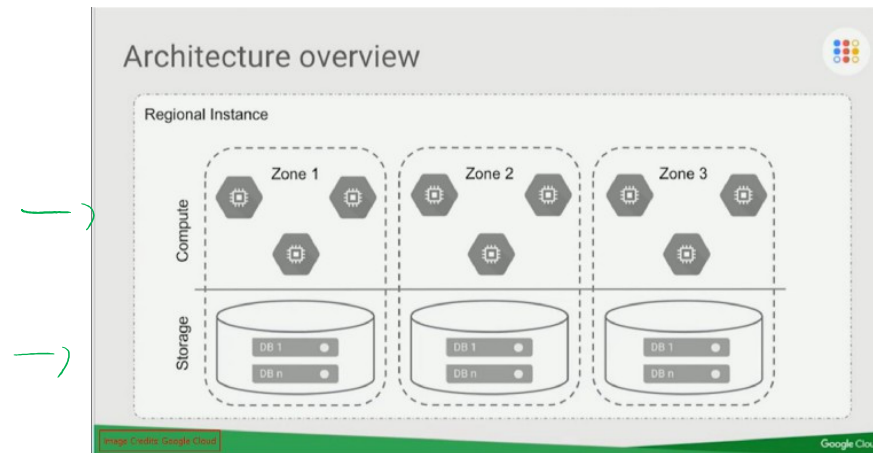


# Prediction

- Best guess is SQL will win

## Prediction

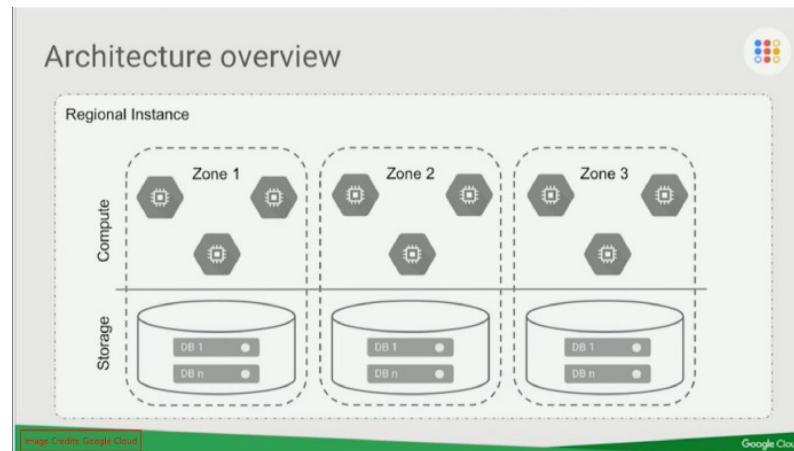
- Best guess is SQL will win
- Pieces are out there already
  - Spanner: multi-node transactions



- AsterixDB: multi-node query optimization

## Prediction

- Best guess is SQL will win
- Pieces are out there already
  - Spanner: multi-node transactions



- AsterixDB: multi-node query optimization
- For now, NoSQL still offers key benefits

INTRO TO DATA STRUCTURE

**JSON**

## Where Are We

- So far, we have studied relational data model
  - Data are stored in tables (relations)
  - Queries are expressions in the [SQL](#)/ [Datalog](#)/ [Relational Algebra](#)



## Where Are We

- So far, we have studied relational data model
  - Data are stored in tables (relations)
  - Queries are expressions in the [SQL](#)/ [Datalog](#)/ [Relational Algebra](#)
- Today: Semi-structured data model
  - Popular formats: XML, [JSon](#), protobuf



## Semi-structured Data

- Database of semi-structured data
  - Collection of nodes(leaf, interior)

## Semi-structured Data

- Database of semi-structured data
  - Collection of nodes(leaf, interior)
- Leaf node
  - Associated data; any atomic type: integer, string, etc





## Semi-structured Data

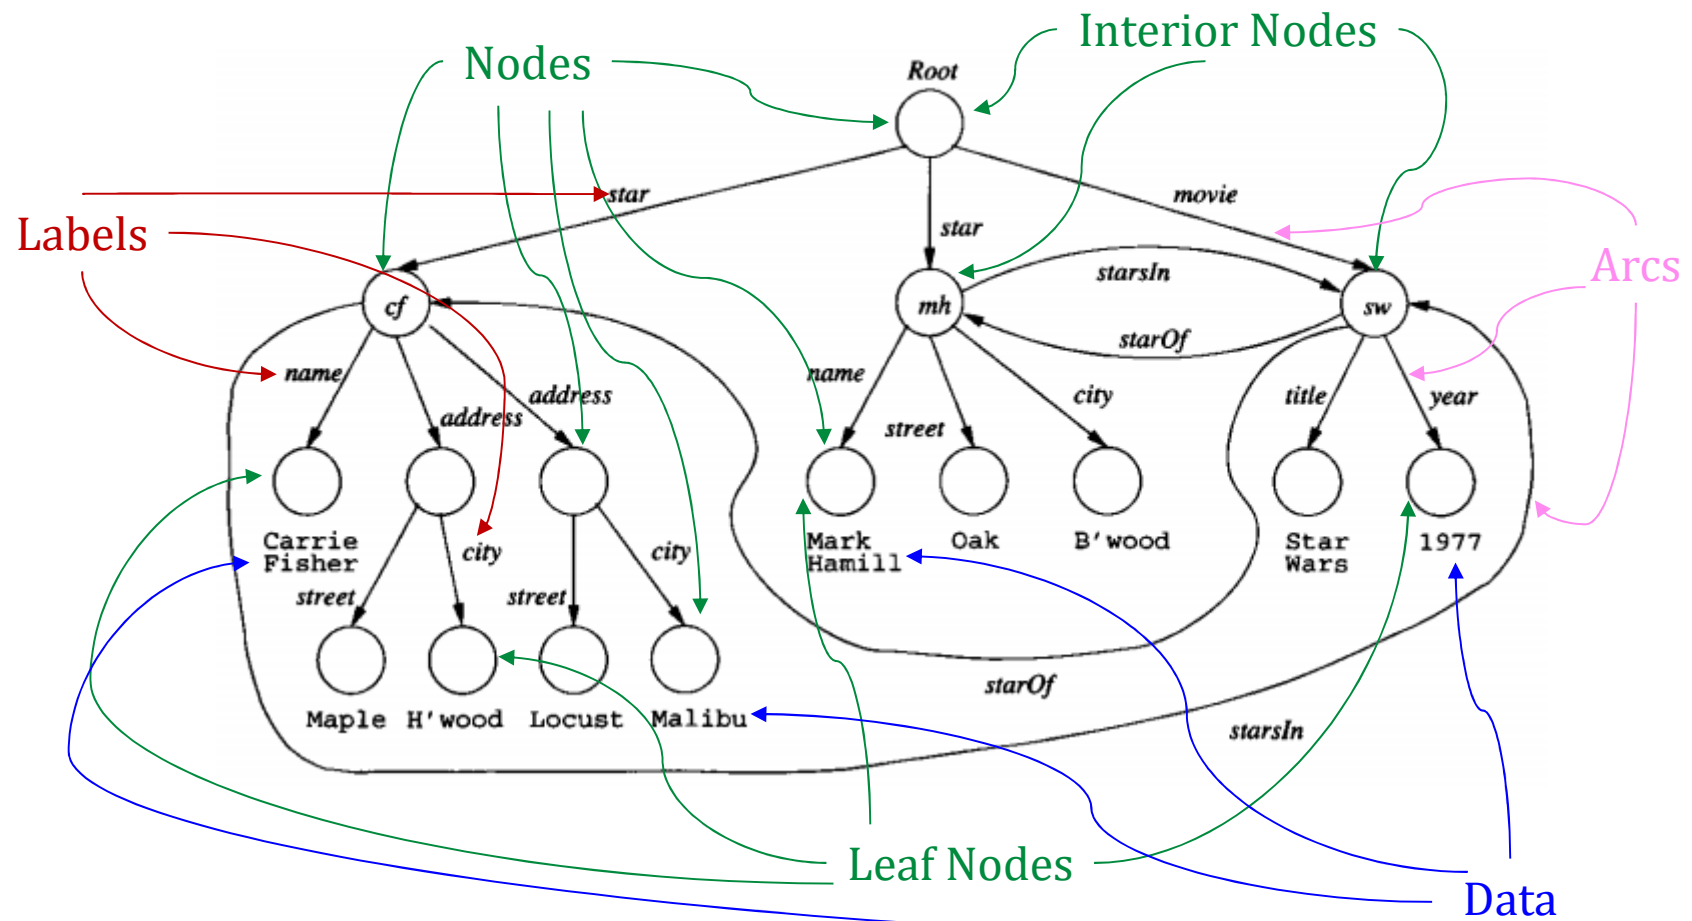
- Database of semi-structured data
  - Collection of nodes(leaf, interior)
- Leaf node
  - Associated data; any atomic type: integer, string, etc
- Interior node
  - One or more arcs going out, each arc has label

## Semi-structured Data

- Database of semi-structured data
  - Collection of nodes(leaf, interior)
- Leaf node
  - Associated data; any atomic type: integer, string, etc
- Interior node
  - One or more arcs going out, each arc has label
- Root node (an interior node)
  - No arc entering, represent entire DB
  - Every node must be reachable from root node



# Semi-structured Data



# JSON

- 13 years ago
  - Java interpreters were **very slow**
  - Native browser function parsed JSON **100x faster**

\* parse = break up a sentence or group of words into separate components, including the definition of each part's function or form



# JSON

- 13 years ago
  - Java interpreters were **very slow**
  - Native browser function parsed JSON **100x faster**
- XML was also an option, but
  - IE had **memory leak** in its XML parser

\* parse = break up a sentence or group of words into separate components, including the definition of each part's function or form



# JSON

- 13 years ago
  - Java interpreters were **very slow**
  - Native browser function parsed JSON **100x faster**
- XML was also an option, but
  - IE had **memory leak** in its XML parser
- JSON used in **Gmail** etc. for this reason
- Spread organically to **server-side** systems

\* parse = break up a sentence or group of words into separate components, including the definition of each part's function or form



## Overview of JSON

- JavaScript Object Notation
  - Lightweight **text-based** open standard designed for human-readable data interchange.
  - Interfaces in C, C++, Java, Python, Perl, etc.

## Overview of JSON

- JavaScript Object Notation
  - Lightweight **text-based** open standard designed for human-readable data interchange.
  - Interfaces in C, C++, Java, Python, Perl, etc.
- The filename extension is **.json**

We will emphasize JSon as semi-structured data





## JSon vs Relational

- Relational data model
  - Rigid flat structure (tables)
  - Schema must be fixed in advanced
  - Binary representation
    - *good for performance, bad for exchange*

## JSon vs Relational

- Relational data model
  - Rigid flat structure (tables)
  - Schema must be fixed in advanced
  - Binary representation
    - *good for performance, bad for exchange*
- Semi-structured data model/JSon
  - Flexible, nested structure (trees)
  - Does not require predefined schema (“self-describing”)
  - Text representation
    - *bad for performance, good for exchange*
  - Most common use
    - *language API; query languages emerging*



## JSon Syntax

```
{ "book": [  
  {"id": "01",  
   "language": "Java",  
   "author": "H. Javeson",  
   "year": 2015  
  },  
  {"id": "07",  
   "language": "C++",  
   "edition": "second",  
   "author": "E. Sepp",  
   "price": 22.25  
  }  
]  
}
```

## JSon Terminology

- Curly braces “{ }” hold objects
  - Each object is a list of **name**/**value** pairs separated by a comma “,”
  - Each pair is a **name** followed by a colon “:”, and followed by the **value**

## JSon Terminology

- Curly braces “{ }” hold objects
  - Each object is a list of **name**/**value** pairs separated by a comma “,”
  - Each pair is a **name** followed by a colon “:”, and followed by the **value**
- Square brackets “[ ]” holds **array** and **values** are separated by comma “,”

## JSon Terminology

- Curly braces “{ }” hold objects
  - Each object is a list of **name**/**value** pairs separated by a comma “,”
  - Each pair is a **name** followed by a colon “:”, and followed by the **value**
- Square brackets “[ ]” holds **array** and **values** are separated by comma “,”
- **Data** made up of **objects**, **lists**, and **atomic values** (**integers**, **floats**, **strings**, **booleans**)



## JSon Data Structures

- Collections of **name-value** pairs:
  - {“**name1**”:**value1**, “**name2**”:**value2**, ...}
  - The “**name**” is also called a “**key**”

## JSon Data Structures

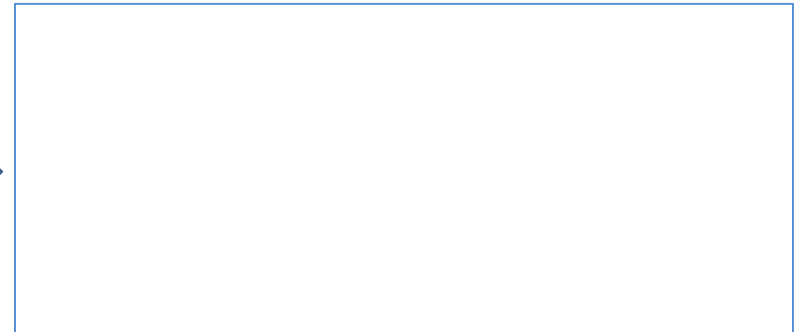

- Collections of **name-value** pairs:
  - {**“name1”**:**value1**, **“name2”**:**value2**, ...}
  - The **“name”** is also called a **“key”**
- Ordered lists of values:
  - [**obj1**, **obj2**, **obj3**, ...]



## Avoid Using Duplicate Keys

The standard allows them, but many implementations doesn't


```
{ "id": "07",  
  "title": "Databases",  
  "author": "Garcia-Molina",  
  "author": "Ullman",  
  "author": "Widom",  
}
```



## Avoid Using Duplicate Keys

The standard allows them, but many implementations doesn't

```
{ "id": "07",  
  "title": "Databases",  
  "author": "Garcia-Molina",  
  "author": "Ullman",  
  "author": "Widom",  
}
```



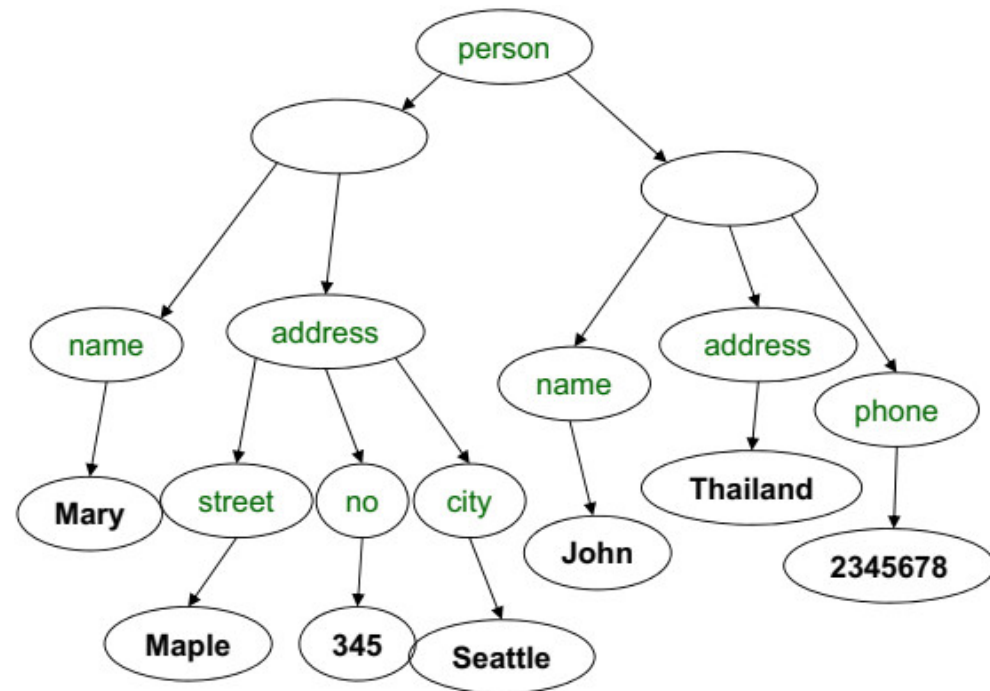
```
{ "id": "07",  
  "title": "Databases",  
  "author": ["Garcia-Molina",  
             "Ullman",  
             "Widom"]  
}
```

## JSon Data Types

- Number
- String = double-quoted
- Boolean = True or False
- Null/empty

# JSon Semantics

```
{ "person":  
  [{ "name": "Mary",  
    "address":  
      { "street": "Maple",  
        "no": 345,  
        "city": "Seattle" } },  
    { "name": "John",  
      "address": "Thailand",  
      "phone": 2345678 }  
  ]  
}
```



## JSon Data

- JSon is self-describing



## JSon Data

- JSon is **self-describing**
- Schema elements become part of the data
  - Relational schema: **person(name, phone)**
  - In JSon: “**person**”, “**name**”, “**phone**” are part of the data, are repeated many times

## JSon Data

- JSon is **self-describing**
- Schema elements become part of the data
  - Relational schema: **person(name, phone)**
  - In JSon: “**person**”, “**name**”, “**phone**” are part of the data, are repeated many times
- Consequence: JSon is much more **flexible**
  - also uses more space (but can be compressed)

## JSon Data

- JSon is **self-describing**
- Schema elements become part of the data
  - Relational schema: **person(name, phone)**
  - In JSon: “**person**”, “**name**”, “**phone**” are part of the data, are repeated many times
- Consequence: JSon is much more **flexible**
  - also uses more space (but can be compressed)
- JSon is an example of **semi-structured** data





**Thank you.**